

Lab 8: Define and Solve an ML Problem of Your Choosing

```
In [18]: import pandas as pd
import numpy as np
import os
import matplotlib.pyplot as plt
import seaborn as sns
```

In this lab assignment, you will follow the machine learning life cycle and implement a model to solve a machine learning problem of your choosing. You will select a data set and choose a predictive problem that the data set supports. You will then inspect the data with your problem in mind and begin to formulate a project plan. You will then implement the machine learning project plan.

You will complete the following tasks:

1. Build Your DataFrame
2. Define Your ML Problem
3. Perform exploratory data analysis to understand your data.
4. Define Your Project Plan
5. Implement Your Project Plan:
 - Prepare your data for your model.
 - Fit your model to the training data and evaluate your model.
 - Improve your model's performance.

Part 1: Build Your DataFrame

You will have the option to choose one of four data sets that you have worked with in this program:

- The "census" data set that contains Census information from 1994:
`censusData.csv`
- Airbnb NYC "listings" data set: `airbnbListingsData.csv`
- World Happiness Report (WHR) data set: `WHR2018Chapter20onlineData.csv`
- Book Review data set: `bookReviewsData.csv`

Note that these are variations of the data sets that you have worked with in this program. For example, some do not include some of the preprocessing necessary for specific models.

Load a Data Set and Save it as a Pandas DataFrame

The code cell below contains filenames (path + filename) for each of the four data sets available to you.

Task: In the code cell below, use the same method you have been using to load the data using `pd.read_csv()` and save it to DataFrame `df`.

You can load each file as a new DataFrame to inspect the data before choosing your data set.

```
In [19]: # File names of the four data sets
adultDataSet_filename = os.path.join(os.getcwd(), "data", "censusData.csv")
airbnbDataSet_filename = os.path.join(os.getcwd(), "data", "airbnbListingsDa
WHRDataSet_filename = os.path.join(os.getcwd(), "data", "WHR2018Chapter20nli
bookReviewDataSet_filename = os.path.join(os.getcwd(), "data", "bookReviewsD

df = pd.read_csv(airbnbDataSet_filename)# YOUR CODE HERE

df
```

Out [19]:

	name	description	neighborhood_overview	host_name	host_location
0	Skylit Midtown Castle	Beautiful, spacious skylit studio in the heart...	Centrally located in the heart of Manhattan ju...	Jennifer	New York, New York, United States
1	Whole flr w/private bdrm, bath & kitchen(pls r...	Enjoy 500 s.f. top floor in 1899 brownstone, w...	Just the right mix of urban center and local n...	LisaRoxanne	New York, New York, United States
2	Spacious Brooklyn Duplex, Patio + Garden	We welcome you to stay in our lovely 2 br dupl...	NaN	Rebecca	Brooklyn, New York, United States
3	Large Furnished Room Near B'way	Please don't expect the luxury here just a bas...	Theater district, many restaurants around here.	Shunichi	New York, New York, United States
4	Cozy Clean Guest Room - Family Apt	Our best guests are seeking a safe, clean, spa...	Our neighborhood is full of restaurants and ca...	MaryEllen	New York, New York, United States
...
28017	Astoria Luxury suite 2A	THIS LOVELY HOME IS THE SPACIOUS SUITE WITH PR...	NaN	Vicky	Queens, New York, United States
28018	Newly renovated suite in the heart of Williams...	Just fully renovated from head to toe. On the ...	NaN	Samuel	New York, New York, United States
28019	Perfect Room to Stay in Brooklyn! Near Metro!	Amazing and comfortable space in Brooklyn, sam...	NaN	Carlos	US
28020	New Beautiful Modern One Bedroom	This stylish place to stay is perfect for a gr...	NaN	Lexia	New York, New York, United States

	name	description	neighborhood_overview	host_name	host_location
	in Brooklyn				
28021	Large, modern, private 1 bedroom in beach condo	Private bedroom on its own floor with very lar...	Beach, surf shop, stop and shop, Dunkin' Donut...	Justine	US

28022 rows × 50 columns

Part 2: Define Your ML Problem

Next you will formulate your ML Problem. In the markdown cell below, answer the following questions:

1. List the data set you have chosen.
2. What will you be predicting? What is the label?
3. Is this a supervised or unsupervised learning problem? Is this a clustering, classification or regression problem? Is it a binary classificaiton or multi-class classifiction problem?
4. What are your features? (note: this list may change after your explore your data)
5. Explain why this is an important problem. In other words, how would a company create value with a model that predicts this label?

Dataset Chosen

Airbnb Listings Data Set

Prediction Task

Prediction: Whether an Airbnb listing has availability. **Label:** `has_availability`

Type of Learning Problem

Learning Type: Supervised learning

Problem Type: Classification

Classification Type: Binary classification (True or False)

Features

Initially considered features:

- `host_response_rate`

- `host_acceptance_rate`
- `host_is_superhost`
- `host_listings_count`
- `review_scores_communication`
- `review_scores_location`
- `review_scores_value`
- `instant_bookable`
- `calculated_host_listings_count`
- `calculated_host_listings_count_entire_homes`
- `calculated_host_listings_count_private_rooms`
- `calculated_host_listings_count_shared_rooms`
- `reviews_per_month`
- `n_host_verifications`

Note: This list may change after data exploration and preprocessing.

Importance of the Problem

Value Creation:

- **Improved Guest Experience:** Ensures guests see available listings, reducing frustration and enhancing satisfaction.
- **Optimized Host Management:** Helps hosts manage listings better, adjusting strategies to maximize occupancy.
- **Platform Efficiency:** Optimizes search results, leading to higher conversion rates and better resource utilization.
- **Revenue Maximization:** Higher occupancy rates mean more revenue for both the platform and hosts.
- **Strategic Insights:** Provides valuable insights into factors affecting availability, helping hosts improve listing performance.

Part 3: Understand Your Data

The next step is to perform exploratory data analysis. Inspect and analyze your data set with your machine learning problem in mind. Consider the following as you inspect your data:

1. What data preparation techniques would you like to use? These data preparation techniques may include:
 - addressing missingness, such as replacing missing values with means
 - finding and replacing outliers
 - renaming features and labels
 - finding and replacing outliers

- performing feature engineering techniques such as one-hot encoding on categorical features
 - selecting appropriate features and removing irrelevant features
 - performing specific data cleaning and preprocessing techniques for an NLP problem
 - addressing class imbalance in your data sample to promote fair AI
2. What machine learning model (or models) you would like to use that is suitable for your predictive problem and data?
- Are there other data preparation techniques that you will need to apply to build a balanced modeling data set for your problem and model? For example, will you need to scale your data?
3. How will you evaluate and improve the model's performance?
- Are there specific evaluation metrics and methods that are appropriate for your model?

Think of the different techniques you have used to inspect and analyze your data in this course. These include using Pandas to apply data filters, using the Pandas `describe()` method to get insight into key statistics for each column, using the Pandas `dtypes` property to inspect the data type of each column, and using Matplotlib and Seaborn to detect outliers and visualize relationships between features and labels. If you are working on a classification problem, use techniques you have learned to determine if there is class imbalance.

Task: Use the techniques you have learned in this course to inspect and analyze your data. You can import additional packages that you have used in this course that you will need to perform this task.

Note: You can add code cells if needed by going to the **Insert** menu and clicking on **Insert Cell Below** in the drop-down menu.

Part 4: Define Your Project Plan

Now that you understand your data, in the markdown cell below, define your plan to implement the remaining phases of the machine learning life cycle (data preparation, modeling, evaluation) to solve your ML problem. Answer the following questions:

- Do you have a new feature list? If so, what are the features that you chose to keep and remove after inspecting the data?
- Explain different data preparation techniques that you will use to prepare your data for modeling.
- What is your model (or models)?

- Describe your plan to train your model, analyze its performance and then improve the model. That is, describe your model building, validation and selection plan to produce a model that generalizes well to new data.

Plan for Implementing Remaining Phases of the Machine Learning Life Cycle

Feature List

After inspecting the data, I decided to keep the following features:

- Numerical: `host_response_rate`, `host_acceptance_rate`, `host_listings_count`, `host_total_listings_count`, `accommodates`, `bathrooms`, `bedrooms`, `beds`, `price`, `minimum_nights`, `maximum_nights`, `availability_30`, `availability_60`, `availability_90`, `availability_365`, `number_of_reviews`, `review_scores_rating`, `review_scores_cleanliness`, `review_scores_checkin`, `review_scores_communication`, `review_scores_location`, `review_scores_value`, `calculated_host_listings_count`, `reviews_per_month`
- One-hot encoded categorical: `host_is_superhost`, `instant_bookable`

Removed features include columns with high missing values, irrelevant columns, and those dropped due to preprocessing.

Data Preparation Techniques

1. **Handling Missing Values:** Fill missing numeric values with the mean and categorical values with the mode.
2. **Convert Percentages:** Convert percentage values to float.
3. **Binary Conversion:** Convert boolean columns to integers.
4. **Standardization:** Scale numerical features using `StandardScaler`.
5. **Remove Constant Features:** Use `VarianceThreshold` to remove features with zero variance.
6. **One-hot Encoding:** One-hot encode binary categorical variables.

Model(s)

1. **Logistic Regression:** For initial model building and evaluation.
2. **Random Forest Classifier:** For improved performance with hyperparameter tuning.

Plan for Model Training, Analysis, and Improvement

1. **Model Building:**
 - **Initial Model:** Train a Logistic Regression model to set a baseline.
 - **Feature Selection:** Use `SelectKBest` to identify the top 10 features.

2. Model Validation:

- **Cross-Validation:** Use 5-fold cross-validation to evaluate model performance.
- **Hyperparameter Tuning:** Use GridSearchCV to find the best hyperparameters for the Random Forest model.

3. Model Analysis:

- **Performance Metrics:** Evaluate using accuracy, precision, recall, F1 score, and ROC AUC score.

4. Model Improvement:

- **Feature Engineering:** Explore additional feature creation or transformation if needed.
- **Ensemble Methods:** Consider combining multiple models if individual model performance is insufficient.

5. Model Selection:

- **Best Model Selection:** Select the model with the best cross-validation performance.
- **Final Evaluation:** Evaluate the final model on a separate test set to ensure generalizability.

This plan will guide the preparation, modeling, and iterative improvement to build a robust model that generalizes well to new data.

Part 5: Implement Your Project Plan

Task: In the code cell below, import additional packages that you have used in this course that you will need to implement your project plan.

```
In [20]: import pandas as pd
import os
from sklearn.preprocessing import StandardScaler
from sklearn.model_selection import train_test_split, GridSearchCV
from sklearn.linear_model import LogisticRegression
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_score
from sklearn.impute import SimpleImputer
from sklearn.compose import ColumnTransformer
from sklearn.pipeline import Pipeline
from sklearn.feature_selection import SelectKBest, f_classif
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.pipeline import make_pipeline
from sklearn.preprocessing import OneHotEncoder
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.metrics import roc_curve, auc, precision_recall_curve
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.metrics import confusion_matrix, roc_curve, auc
```


Task: Use the rest of this notebook to carry out your project plan.

You will:

1. Prepare your data for your model.
2. Fit your model to the training data and evaluate your model.
3. Improve your model's performance by performing model selection and/or feature selection techniques to find best model for your problem.

Add code cells below and populate the notebook with commentary, code, analyses, results, and figures as you see fit.

```
In [21]: # Import necessary libraries
import pandas as pd
import numpy as np
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler, OneHotEncoder
from sklearn.compose import ColumnTransformer
from sklearn.pipeline import Pipeline
from sklearn.impute import SimpleImputer
from sklearn.linear_model import LogisticRegression
from sklearn.ensemble import RandomForestClassifier
from sklearn.model_selection import GridSearchCV
from sklearn.feature_selection import SelectKBest, f_classif, VarianceThreshold
from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_score

# Load the Airbnb dataset into a pandas DataFrame
airbnb_df = df # Use the correct path to the dataset

# Drop unnecessary columns
airbnb_df.drop(columns=[
    'name', 'room_type', 'neighbourhood_group_cleansed', 'amenities', 'host_description', 'neighborhood_overview', 'host_name', 'host_about'
], inplace=True)

# Display the first few rows of the DataFrame to confirm successful loading
print("Airbnb Listings Data Set:")
print(airbnb_df.head())

# Check for missing values
print(airbnb_df.isnull().sum())

# Get a summary of the dataset
print(airbnb_df.info())
print(airbnb_df.describe())

# Fill missing numeric values with the mean
for col in ['host_response_rate', 'host_acceptance_rate', 'reviews_per_month']:
    airbnb_df[col].fillna(airbnb_df[col].mean(), inplace=True)

# Drop rows with missing target values
airbnb_df.dropna(subset=['has_availability'], inplace=True)
```

```

# Fill other missing categorical values with the mode
for col in ['host_is_superhost', 'instant_bookable']:
    airbnb_df[col].fillna(airbnb_df[col].mode()[0], inplace=True)

def convert_percentage_to_float(x):
    return float(x.strip('%')) / 100 if isinstance(x, str) else x

airbnb_df['host_response_rate'] = airbnb_df['host_response_rate'].apply(convert_percentage_to_float)
airbnb_df['host_acceptance_rate'] = airbnb_df['host_acceptance_rate'].apply(convert_percentage_to_float)

# Convert 'has_availability' to binary (True -> 1, False -> 0)
airbnb_df['has_availability'] = airbnb_df['has_availability'].apply(lambda x: 1 if x else 0)

# Convert boolean columns to integers
bool_cols = ['host_has_profile_pic', 'host_identity_verified', 'instant_bookable']
for col in bool_cols:
    airbnb_df[col] = airbnb_df[col].astype(int)

# One-hot encode binary categorical variables
airbnb_df = pd.get_dummies(airbnb_df, columns=['host_is_superhost', 'instant_bookable'])

# List of numerical features to scale
num_features = [
    'host_response_rate', 'host_acceptance_rate', 'host_listings_count',
    'host_total_listings_count', 'host_has_profile_pic',
    'host_identity_verified', 'accommodates', 'bathrooms', 'bedrooms',
    'beds', 'price', 'minimum_nights', 'maximum_nights',
    'minimum_minimum_nights', 'maximum_minimum_nights',
    'minimum_maximum_nights', 'maximum_maximum_nights',
    'minimum_nights_avg_ntm', 'maximum_nights_avg_ntm',
    'availability_30', 'availability_60', 'availability_90',
    'availability_365', 'number_of_reviews', 'number_of_reviews_ltm',
    'number_of_reviews_l30d', 'review_scores_rating',
    'review_scores_cleanliness', 'review_scores_checkin',
    'review_scores_communication', 'review_scores_location',
    'review_scores_value', 'calculated_host_listings_count',
    'calculated_host_listings_count_entire_homes',
    'calculated_host_listings_count_private_rooms',
    'calculated_host_listings_count_shared_rooms', 'reviews_per_month',
    'n_host_verifications'
]

# Remove the target column from the numerical features list
if 'has_availability' in num_features:
    num_features.remove('has_availability')

# Verify numerical features exist in DataFrame
for feature in num_features:
    if feature not in airbnb_df.columns:
        print(f"Warning: Feature '{feature}' not found in DataFrame")

# Standardize the numerical features
scaler = StandardScaler()
airbnb_df[num_features] = scaler.fit_transform(airbnb_df[num_features])

# Check if all columns are numeric

```

```

print(airbnb_df.dtypes)

# Check the columns in the DataFrame
print("DataFrame columns:", airbnb_df.columns)

# Define numerical and categorical columns
numerical_cols = airbnb_df.select_dtypes(include=['int64', 'float64']).columns
categorical_cols = airbnb_df.select_dtypes(include=['object', 'bool', 'uint8']).columns

# Define the target column
target_column = 'has_availability'

# Remove the target column from numerical and categorical features list if it is present
if target_column in numerical_cols:
    numerical_cols.remove(target_column)
if target_column in categorical_cols:
    categorical_cols.remove(target_column)

# Verify columns before preprocessing
print("Numerical columns before preprocessing:", numerical_cols)
print("Categorical columns before preprocessing:", categorical_cols)

# Define preprocessing steps
preprocessor = ColumnTransformer(
    transformers=[
        ('num', Pipeline(steps=[
            ('imputer', SimpleImputer(strategy='mean')),
            ('scaler', StandardScaler())
        ]), numerical_cols),
        ('cat', Pipeline(steps=[
            ('imputer', SimpleImputer(strategy='most_frequent')),
            ('onehot', OneHotEncoder(handle_unknown='ignore'))
        ]), categorical_cols)
    ]
)

# Define features and target variable
X = airbnb_df.drop('has_availability', axis=1)
y = airbnb_df['has_availability']

# Verify columns before transformation
print("Columns in X before transformation:", X.columns)

try:
    X_transformed = preprocessor.fit_transform(X)
except ValueError as e:
    missing_columns = set(numerical_cols + categorical_cols) - set(X.columns)
    raise ValueError(f"The following columns are missing in the DataFrame: {missing_columns}")

# Verify there are no NaN values in the transformed data
if pd.DataFrame(X_transformed).isnull().sum().sum() > 0:
    raise ValueError("There are still missing values in the transformed features")

# Verify transformation output
print("Initial X_transformed shape:", X_transformed.shape)

```

```

# Remove constant features
constant_filter = VarianceThreshold(threshold=0)
X_transformed = constant_filter.fit_transform(X_transformed)

# Verify shape after removing constant features
print("Shape after VarianceThreshold:", X_transformed.shape)

# Select top k features
k = 10
selector = SelectKBest(score_func=f_classif, k=k)

if pd.DataFrame(X_transformed).isnull().sum().sum() > 0:
    raise ValueError("There are still missing values in the transformed feat

X_new = selector.fit_transform(X_transformed, y)

# Verify shape after feature selection
print("Shape after SelectKBest:", X_new.shape)

# Split the data into training and testing sets with selected features
X_train_new, X_test_new, y_train_new, y_test_new = train_test_split(X_new, y

# Verify shapes after train-test split
print("Shapes after train-test split - X_train_new:", X_train_new.shape, "X_

# Define the model
rf = RandomForestClassifier(random_state=42)

# Define the parameter grid
param_grid = {
    'n_estimators': [50, 100, 200],
    'max_depth': [None, 10, 20, 30],
    'min_samples_split': [2, 5, 10],
    'min_samples_leaf': [1, 2, 4]
}

# Debug statement to confirm we reached here
print("Starting GridSearchCV")

# Perform grid search
try:
    grid_search = GridSearchCV(estimator=rf, param_grid=param_grid, cv=5, sc
    grid_search.fit(X_train_new, y_train_new)
    print("GridSearchCV completed")
except Exception as e:
    print("Error during GridSearchCV:", e)
    raise

# Get the best model
best_rf = grid_search.best_estimator_

# Make predictions with the best model
y_pred_new = best_rf.predict(X_test_new)

# Evaluate the best model
print("Random Forest Classifier with Grid Search:")

```

```
print("Accuracy:", accuracy_score(y_test_new, y_pred_new))  
print("Precision:", precision_score(y_test_new, y_pred_new))  
print("Recall:", recall_score(y_test_new, y_pred_new))  
print("F1 Score:", f1_score(y_test_new, y_pred_new))  
print("ROC AUC Score:", roc_auc_score(y_test_new, y_pred_new))
```

Airbnb Listings Data Set:

	host_response_rate	host_acceptance_rate	host_is_superhost	\
0	0.80	0.17	True	
1	0.09	0.69	True	
2	1.00	0.25	True	
3	1.00	1.00	True	
4	NaN	NaN	True	

	host_listings_count	host_total_listings_count	host_has_profile_pic	\
0	8.0	8.0	True	
1	1.0	1.0	True	
2	1.0	1.0	True	
3	1.0	1.0	True	
4	1.0	1.0	True	

	host_identity_verified	accommodates	bathrooms	bedrooms	...	\
0	True	1	1.0	NaN	...	
1	True	3	1.0	1.0	...	
2	True	4	1.5	2.0	...	
3	True	2	1.0	1.0	...	
4	True	1	1.0	1.0	...	

	review_scores_communication	review_scores_location	review_scores_value	\
0	4.79	4.86	4.41	
1	4.80	4.71	4.64	
2	5.00	4.50	5.00	
3	4.42	4.87	4.36	
4	4.95	4.94	4.92	

	instant_bookable	calculated_host_listings_count	\
0	False	3	
1	False	1	
2	False	1	
3	False	1	
4	False	1	

	calculated_host_listings_count_entire_homes	\
0	3	
1	1	
2	1	
3	0	
4	0	

	calculated_host_listings_count_private_rooms	\
0	0	
1	0	
2	0	
3	1	
4	1	

	calculated_host_listings_count_shared_rooms	reviews_per_month	\
0	0	0.33	
1	0	4.86	
2	0	0.02	
3	0	3.68	

```
4                                0                                0.87

    n_host_verifications
0                               9
1                               6
2                               3
3                               4
4                               7

[5 rows x 41 columns]
host_response_rate            11843
host_acceptance_rate          11113
host_is_superhost              0
host_listings_count            0
host_total_listings_count      0
host_has_profile_pic           0
host_identity_verified         0
accommodates                   0
bathrooms                      0
bedrooms                       2918
beds                           1354
price                          0
minimum_nights                 0
maximum_nights                 0
minimum_minimum_nights         0
maximum_minimum_nights         0
minimum_maximum_nights         0
maximum_maximum_nights         0
minimum_nights_avg_ntm         0
maximum_nights_avg_ntm         0
has_availability               0
availability_30                 0
availability_60                 0
availability_90                 0
availability_365                0
number_of_reviews              0
number_of_reviews_ltm           0
number_of_reviews_l30d          0
review_scores_rating            0
review_scores_cleanliness       0
review_scores_checkin           0
review_scores_communication     0
review_scores_location          0
review_scores_value             0
instant_bookable                0
calculated_host_listings_count  0
calculated_host_listings_count_entire_homes 0
calculated_host_listings_count_private_rooms 0
calculated_host_listings_count_shared_rooms 0
reviews_per_month               0
n_host_verifications           0
dtype: int64
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 28022 entries, 0 to 28021
Data columns (total 41 columns):
#    Column                                Non-Null Count  Dtype
```

```

-----
0  host_response_rate          16179 non-null float64
1  host_acceptance_rate       16909 non-null float64
2  host_is_superhost          28022 non-null bool
3  host_listings_count         28022 non-null float64
4  host_total_listings_count   28022 non-null float64
5  host_has_profile_pic        28022 non-null bool
6  host_identity_verified      28022 non-null bool
7  accommodates                28022 non-null int64
8  bathrooms                   28022 non-null float64
9  bedrooms                    25104 non-null float64
10 beds                        26668 non-null float64
11 price                       28022 non-null float64
12 minimum_nights              28022 non-null int64
13 maximum_nights              28022 non-null int64
14 minimum_minimum_nights      28022 non-null float64
15 maximum_minimum_nights      28022 non-null float64
16 minimum_maximum_nights      28022 non-null float64
17 maximum_maximum_nights      28022 non-null float64
18 minimum_nights_avg_ntm      28022 non-null float64
19 maximum_nights_avg_ntm      28022 non-null float64
20 has_availability             28022 non-null bool
21 availability_30              28022 non-null int64
22 availability_60              28022 non-null int64
23 availability_90              28022 non-null int64
24 availability_365             28022 non-null int64
25 number_of_reviews            28022 non-null int64
26 number_of_reviews_ltm        28022 non-null int64
27 number_of_reviews_l30d       28022 non-null int64
28 review_scores_rating         28022 non-null float64
29 review_scores_cleanliness    28022 non-null float64
30 review_scores_checkin        28022 non-null float64
31 review_scores_communication  28022 non-null float64
32 review_scores_location       28022 non-null float64
33 review_scores_value          28022 non-null float64
34 instant_bookable             28022 non-null bool
35 calculated_host_listings_count 28022 non-null int64
36 calculated_host_listings_count_entire_homes 28022 non-null int64
37 calculated_host_listings_count_private_rooms 28022 non-null int64
38 calculated_host_listings_count_shared_rooms 28022 non-null int64
39 reviews_per_month            28022 non-null float64
40 n_host_verifications         28022 non-null int64

```

dtypes: bool(5), float64(21), int64(15)

memory usage: 7.8 MB

None

	host_response_rate	host_acceptance_rate	host_listings_count \
count	16179.000000	16909.000000	28022.000000
mean	0.906901	0.791953	14.554778
std	0.227282	0.276732	120.721287
min	0.000000	0.000000	0.000000
25%	0.940000	0.680000	1.000000
50%	1.000000	0.910000	1.000000
75%	1.000000	1.000000	3.000000
max	1.000000	1.000000	3387.000000

	host_total_listings_count	accommodates	bathrooms	bedrooms
--	---------------------------	--------------	-----------	----------

\					
count		28022.000000	28022.000000	28022.000000	25104.000000
mean		14.554778	2.874491	1.142174	1.329708
std		120.721287	1.860251	0.421132	0.700726
min		0.000000	1.000000	0.000000	1.000000
25%		1.000000	2.000000	1.000000	1.000000
50%		1.000000	2.000000	1.000000	1.000000
75%		3.000000	4.000000	1.000000	1.000000
max		3387.000000	16.000000	8.000000	12.000000

	beds	price	minimum_nights	...	review_scores_checki
n \					
count	26668.000000	28022.000000	28022.000000	...	28022.000000
0					
mean	1.629556	154.228749	18.689387	...	4.81430
0					
std	1.097104	140.816605	25.569151	...	0.43860
3					
min	1.000000	29.000000	1.000000	...	0.00000
0					
25%	1.000000	70.000000	2.000000	...	4.81000
0					
50%	1.000000	115.000000	30.000000	...	4.96000
0					
75%	2.000000	180.000000	30.000000	...	5.00000
0					
max	21.000000	1000.000000	1250.000000	...	5.00000
0					

	review_scores_communication	review_scores_location	\
count	28022.000000	28022.000000	
mean	4.808041	4.750393	
std	0.464585	0.415717	
min	0.000000	0.000000	
25%	4.810000	4.670000	
50%	4.970000	4.880000	
75%	5.000000	5.000000	
max	5.000000	5.000000	

	review_scores_value	calculated_host_listings_count	\
count	28022.000000	28022.000000	
mean	4.647670	9.581900	
std	0.518023	32.227523	
min	0.000000	1.000000	
25%	4.550000	1.000000	
50%	4.780000	1.000000	
75%	5.000000	3.000000	
max	5.000000	421.000000	

	calculated_host_listings_count_entire_homes	\
count	28022.000000	
mean	5.562986	
std	26.121426	
min	0.000000	
25%	0.000000	
50%	1.000000	

75%	1.000000
max	308.000000

	calculated_host_listings_count_private_rooms \
count	28022.000000
mean	3.902077
std	17.972386
min	0.000000
25%	0.000000
50%	0.000000
75%	1.000000
max	359.000000

	calculated_host_listings_count_shared_rooms	reviews_per_month \
count	28022.000000	28022.000000
mean	0.048283	1.758325
std	0.442459	4.446143
min	0.000000	0.010000
25%	0.000000	0.130000
50%	0.000000	0.510000
75%	0.000000	1.830000
max	8.000000	141.000000

	n_host_verifications
count	28022.000000
mean	5.169510
std	2.028497
min	1.000000
25%	4.000000
50%	5.000000
75%	7.000000
max	13.000000

[8 rows x 36 columns]

host_response_rate	float64
host_acceptance_rate	float64
host_listings_count	float64
host_total_listings_count	float64
host_has_profile_pic	float64
host_identity_verified	float64
accommodates	float64
bathrooms	float64
bedrooms	float64
beds	float64
price	float64
minimum_nights	float64
maximum_nights	float64
minimum_minimum_nights	float64
maximum_minimum_nights	float64
minimum_maximum_nights	float64
maximum_maximum_nights	float64
minimum_nights_avg_ntm	float64
maximum_nights_avg_ntm	float64
has_availability	int64
availability_30	float64
availability_60	float64

```

availability_90                float64
availability_365               float64
number_of_reviews              float64
number_of_reviews_ltm          float64
number_of_reviews_l30d         float64
review_scores_rating            float64
review_scores_cleanliness       float64
review_scores_checkin           float64
review_scores_communication     float64
review_scores_location          float64
review_scores_value             float64
calculated_host_listings_count float64
calculated_host_listings_count_entire_homes float64
calculated_host_listings_count_private_rooms float64
calculated_host_listings_count_shared_rooms float64
reviews_per_month              float64
n_host_verifications           float64
host_is_superhost_True          uint8
instant_bookable_0              uint8
instant_bookable_1              uint8
dtype: object
DataFrame columns: Index(['host_response_rate', 'host_acceptance_rate', 'host_listings_count',
                           'host_total_listings_count', 'host_has_profile_pic',
                           'host_identity_verified', 'accommodates', 'bathrooms', 'bedrooms',
                           'beds', 'price', 'minimum_nights', 'maximum_nights',
                           'minimum_minimum_nights', 'maximum_minimum_nights',
                           'minimum_maximum_nights', 'maximum_maximum_nights',
                           'minimum_nights_avg_ntm', 'maximum_nights_avg_ntm', 'has_availability',
                           'availability_30', 'availability_60', 'availability_90',
                           'availability_365', 'number_of_reviews', 'number_of_reviews_ltm',
                           'number_of_reviews_l30d', 'review_scores_rating',
                           'review_scores_cleanliness', 'review_scores_checkin',
                           'review_scores_communication', 'review_scores_location',
                           'review_scores_value', 'calculated_host_listings_count',
                           'calculated_host_listings_count_entire_homes',
                           'calculated_host_listings_count_private_rooms',
                           'calculated_host_listings_count_shared_rooms', 'reviews_per_month',
                           'n_host_verifications', 'host_is_superhost_True', 'instant_bookable_0',
                           'instant_bookable_1'],
                           dtype='object')

```

Numerical columns before preprocessing: ['host_response_rate', 'host_acceptance_rate', 'host_listings_count', 'host_total_listings_count', 'host_has_profile_pic', 'host_identity_verified', 'accommodates', 'bathrooms', 'bedrooms', 'beds', 'price', 'minimum_nights', 'maximum_nights', 'minimum_minimum_nights', 'maximum_minimum_nights', 'minimum_maximum_nights', 'maximum_maximum_nights', 'minimum_nights_avg_ntm', 'maximum_nights_avg_ntm', 'availability_30', 'availability_60', 'availability_90', 'availability_365', 'number_of_reviews', 'number_of_reviews_ltm', 'number_of_reviews_l30d', 'review_scores_rating', 'review_scores_cleanliness', 'review_scores_checkin', 'review_scores_communication', 'review_scores_location', 'review_scores_value', 'calculated_host_listings_count', 'calculated_host_listings_count_entire_homes', 'calculated_host_listings_count_private_rooms', 'calculated_host_listings_count_shared_rooms', 'reviews_per_month', 'n_host_verifications']

Categorical columns before preprocessing: ['host_is_superhost_True', 'instant_bookable_0', 'instant_bookable_1']

Columns in X before transformation: Index(['host_response_rate', 'host_acceptance_rate', 'host_listings_count',

```

    'host_total_listings_count', 'host_has_profile_pic',
    'host_identity_verified', 'accommodates', 'bathrooms', 'bedrooms',
    'beds', 'price', 'minimum_nights', 'maximum_nights',
    'minimum_minimum_nights', 'maximum_minimum_nights',
    'minimum_maximum_nights', 'maximum_maximum_nights',
    'minimum_nights_avg_ntm', 'maximum_nights_avg_ntm', 'availability_30',
    'availability_60', 'availability_90', 'availability_365',
    'number_of_reviews', 'number_of_reviews_ltm', 'number_of_reviews_l30d',
    'review_scores_rating', 'review_scores_cleanliness',
    'review_scores_checkin', 'review_scores_communication',
    'review_scores_location', 'review_scores_value',
    'calculated_host_listings_count',
    'calculated_host_listings_count_entire_homes',
    'calculated_host_listings_count_private_rooms',
    'calculated_host_listings_count_shared_rooms', 'reviews_per_month',
    'n_host_verifications', 'host_is_superhost_True', 'instant_bookable_0',
    'instant_bookable_1'],
    dtype='object')

```

Initial X_transformed shape: (28022, 43)

Shape after VarianceThreshold: (28022, 40)

Shape after SelectKBest: (28022, 10)

Shapes after train-test split - X_train_new: (22417, 10) X_test_new: (5605, 10) y_train_new: (22417,) y_test_new: (5605,)

Starting GridSearchCV

GridSearchCV completed

Random Forest Classifier with Grid Search:

Accuracy: 0.960392506690455

Precision: 0.960392506690455

Recall: 1.0

F1 Score: 0.9797961412449946

ROC AUC Score: 0.5

```

In [22]: # Plot confusion matrix
conf_matrix = confusion_matrix(y_test_new, y_pred_new)
plt.figure(figsize=(8, 6))
sns.heatmap(conf_matrix, annot=True, fmt='d', cmap='Blues')
plt.title('Confusion Matrix')
plt.xlabel('Predicted')
plt.ylabel('Actual')
plt.show()

# Plot ROC curve
fpr, tpr, thresholds = roc_curve(y_test_new, grid_search.predict_proba(X_test_new)[:,1])
roc_auc = auc(fpr, tpr)

plt.figure(figsize=(8, 6))
plt.plot(fpr, tpr, color='darkorange', lw=2, label=f'ROC curve (area = {roc_auc})')
plt.plot([0, 1], [0, 1], color='navy', lw=2, linestyle='--')
plt.xlim([0.0, 1.0])

```


```



plt.ylim([0.0, 1.05])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver Operating Characteristic (ROC) Curve')
plt.legend(loc='lower right')
plt.show()

# Plot feature importance
feature_importances = best_rf.feature_importances_
features = selector.get_feature_names_out()

plt.figure(figsize=(12, 8))
plt.barh(range(len(feature_importances)), feature_importances, align='center')
plt.yticks(range(len(feature_importances)), features)
plt.xlabel('Feature Importance')
plt.title('Feature Importance from Random Forest')
plt.show()

```





In []:

